

Edge Computing Lab

Class: TY-AIEC

School of Computing, MIT Art Design Technology University

Academic Year: 2024-25

Experiment No. 8

Name: HARSHIT PRAKASH SHETTY

Class: TY-AIEC-B

Enrollment No: MITU22BTCS0321

Roll No: 2223617

Introduction

The "magic wand" project that can recognize gestures using an accelerometer and an ML classification model on Edge Devices

Objective: Build a project to detect the accelerometer values and convert them into gestures

Tasks:

- Generate the dataset for Accelerometer Motion (Up-Down, Left-Right)
- Configure BLE Sense / Mobile for Edge Impulse
- Building and Training a Model
- Deploy on Nano BLE Sense / Mobile Phone

Introduction

Edge Impulse is a development platform for machine learning on edge devices, targeted at developers who want to create intelligent device solutions. The " Accelerometer Motion "sensor reading equivalent in Edge Impulse would typically involve creating a simple machine learning model that can run on an edge device, like classifying sensor data or recognizing a basic pattern.

Materials Required

- Nano BLE Sense Board

Theory

GPIO (General Purpose Input/Output) pins on the Raspberry Pi are used for interfacing with other electronic components. BCM numbering refers to the pin numbers in the Broadcom SOC channel, which is a more consistent way to refer to the GPIO pins across different versions of the

Here's a high-level overview of steps you'd follow to create a "Hello World" project on Edge Impulse:

Steps to Configure the Edge Impulse:

1. Create an Account and New Project:

- Sign up for an Edge Impulse account.
 - Create a new project from the dashboard.
2. Connect a Device:
 - You can use a supported development board or your smartphone as a sensor device.
 - Follow the instructions to connect your device to your Edge Impulse project.
 3. Collect Data:
 - Use the Edge Impulse mobile app or the Web interface to collect data from the onboard sensors.
 - For a "Hello World" project, you could collect accelerometer data, for instance.
 4. Create an Impulse:
 - Go to the 'Create impulse' page.
 - Add a processing block (e.g., time-series data) and a learning block (e.g., classification).
 - Save the impulse, which defines the machine learning pipeline.
 5. Design a Neural Network:
 - Navigate to the 'NN Classifier' under the 'Learning blocks'.
 - Design a simple neural network. Edge Impulse provides a default architecture that works well for most basic tasks.
 6. Train the Model:
 - ☐ Click on the 'Start training' button to train your machine learning model with the collected data.
 7. Test the Model:
 - ☐ Once the model is trained, you can test its performance with new data in the 'Model Testing' tab.
 8. Deploy the Model:

- Go to the 'Deployment' tab.
- Select the deployment method that suits your edge device (e.g., Arduino library, WebAssembly, container, etc.).
- Follow the instructions to deploy the model to your device.

9. Run Inference:

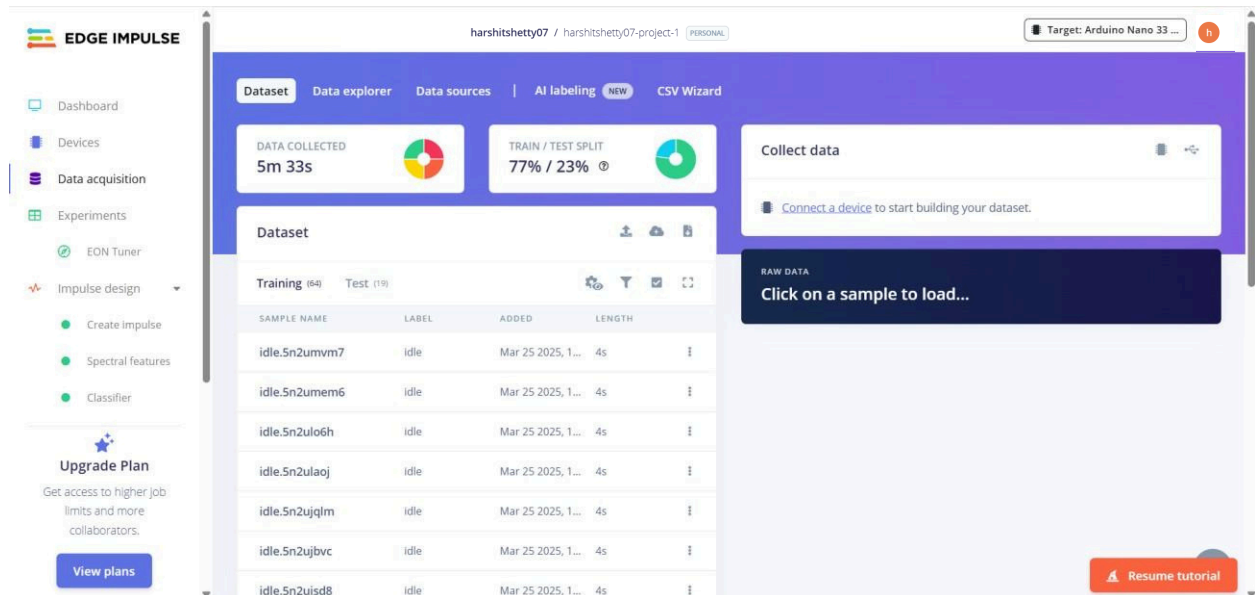
- ☐ With the model deployed, run inference on the edge device to see it classifying data in real-time.

10. Monitor:

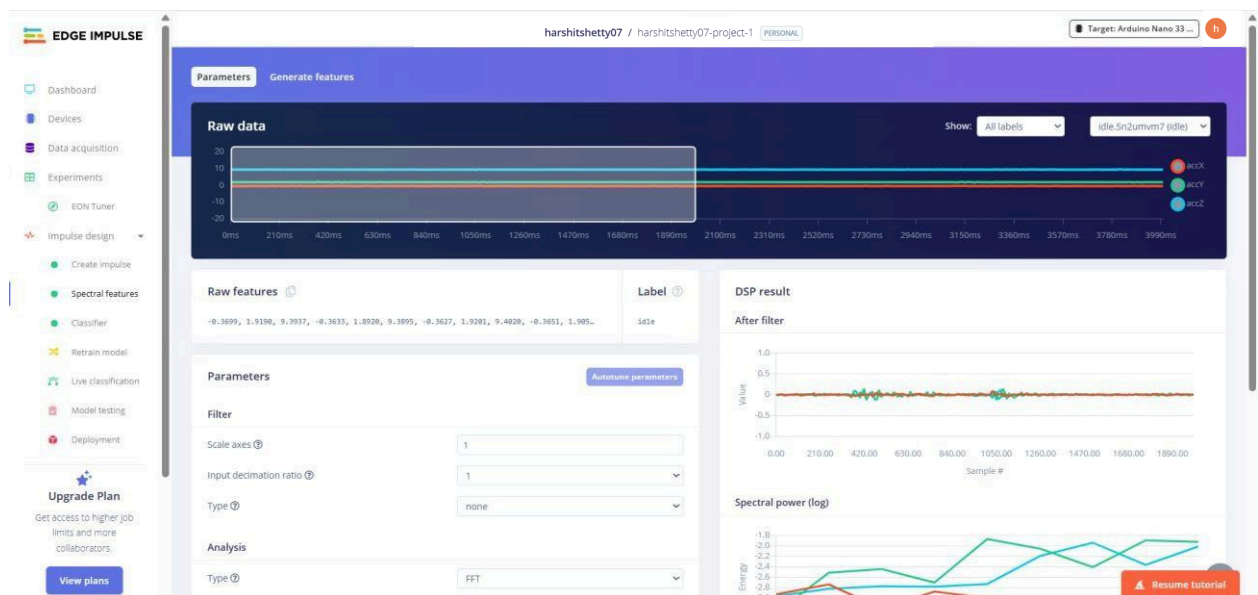
- ☐ You can monitor the performance of your device through the Edge Impulse studio.

Paste your Edge Impulse project's Results:

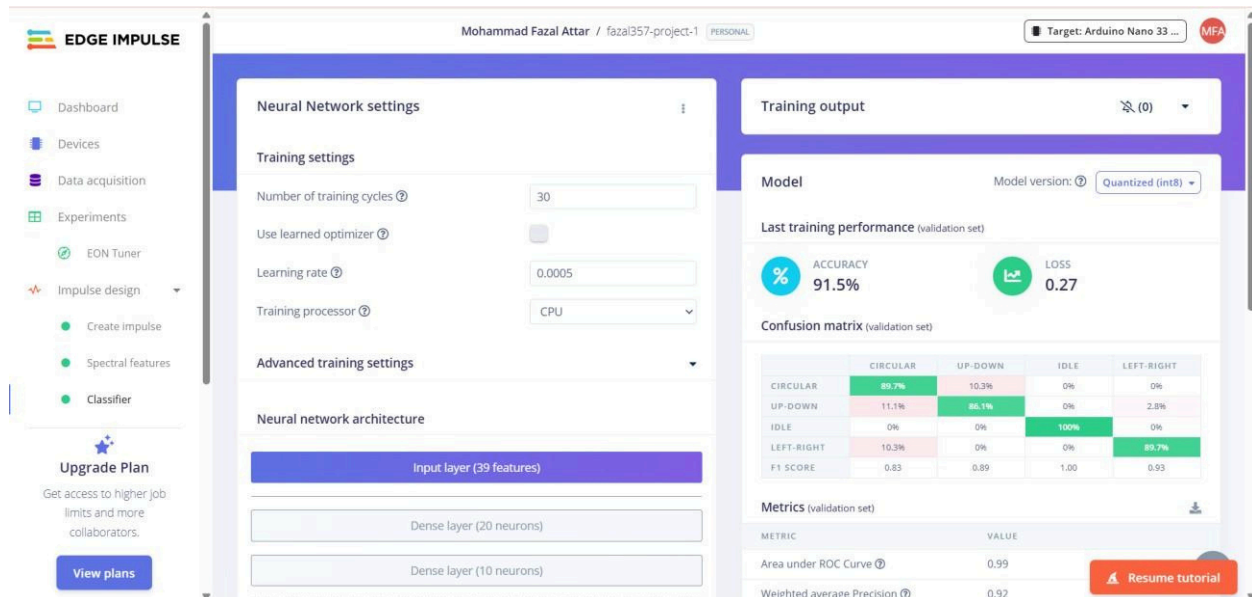
1) Dataset Image



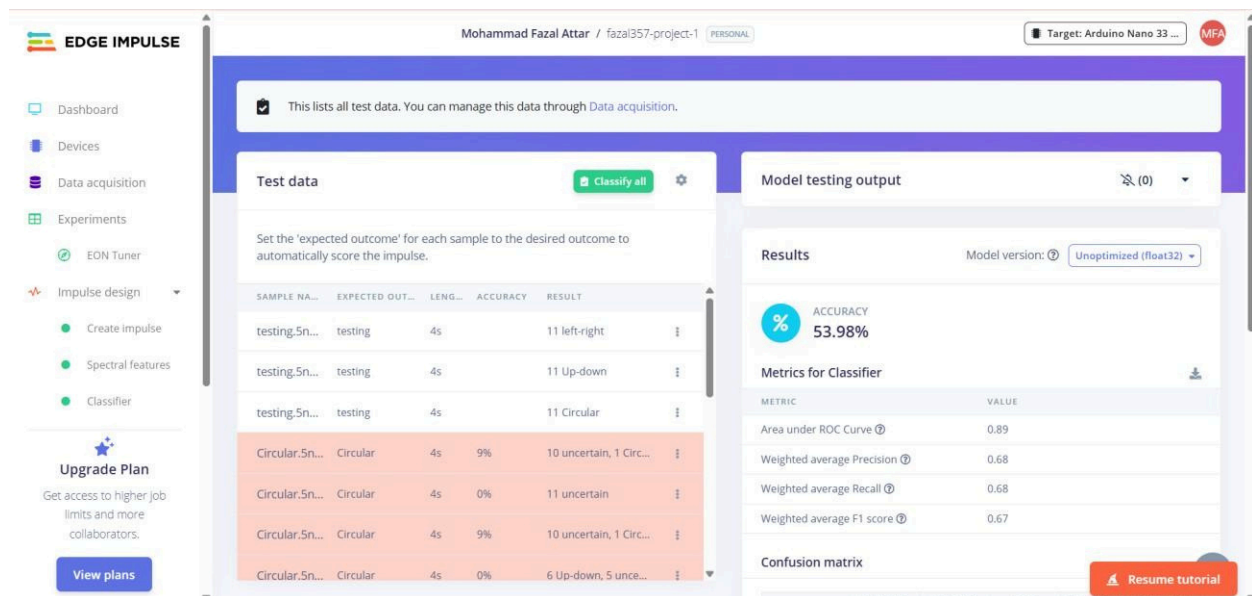
2) Feature extraction – Image



3) Accuracy / Loss - Confusion Matrix – image



4) Validation Result – Image



5) Copy the code of Arduino Sketch

```
/* Edge Impulse ingestion SDK
 * Copyright (c) 2022 EdgeImpulse Inc.
 *
 * Licensed under the Apache License, Version 2.0 (the "License"); * you may
not use this file except in compliance with the License.
 * You may obtain a copy of the License at
 * http://www.apache.org/licenses/LICENSE-2.0
 *
 * Unless required by applicable law or agreed to in writing, software
 * distributed under the License is distributed on an "AS IS" BASIS,
 * WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
 * See the License for the specific language governing permissions and
 * limitations under the License.
 */

/* Includes ----- */
#include <fazal357-project-1_inferencing.h>
#include <Arduino_LSM9DS1.h> //Click here to get the library:
https://www.arduino.cc/reference/en/libraries/arduino_lsm9ds1/

/* Constant defines ----- */
#define CONVERT_G_TO_MS2    9.80665f
/**
 * When data is collected by the Edge Impulse Arduino Nano 33 BLE Sense *
firmware, it is limited to a 2G range. If the model was created with a *
different sample range, modify this constant to match the input values. * See
https://github.com/edgeimpulse/firmware-arduino-nano-33-
blesense/blob/master/src/sensors/ei_lsm9ds1.cpp
 * for more information.
 */
#define MAX_ACCEPTED_RANGE 2.0f
// -DEI_CLASSIFIER_ALLOCATION_STATIC

/*
** NOTE: If you run into TFLite arena allocation issue.
**
** This may be due to may dynamic memory fragmentation.
```

```

** Try defining "-DEI_CLASSIFIER_ALLOCATION_STATIC" in boards.local.txt (create
** if it doesn't exist) and copy this file to
** `https://support.arduino.cc/hc/en-us/articles/360012076960-Where-are-
theinstalled-cores-located-)
** to find where Arduino installs cores on your machine.
**
** If the problem persists then there's not enough memory for this model and
application.
*/

/* Private variables ----- */
static bool debug_nn = false; // Set this to true to see e.g. features generated
from the raw signal
static uint32_t run_inference_every_ms = 200; static
rtos::Thread inference_thread(osPriorityLow); static float
buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE] = { 0 }; static
float inference_buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE];

/* Forward declaration */
void
run_inference_background();

/**
 * @brief Arduino setup function
 */ void
setup()
{
    // put your setup code here, to run once:
    Serial.begin(115200);
    // comment out the below line to cancel the wait for USB connection
    (needed for native USB) while (!Serial);
    Serial.println("Edge Impulse Inferencing Demo");
    if (!IMU.begin()) {
ei_printf("Failed to initialize IMU!\r\n");
    } else {
ei_printf("IMU initialized\r\n");
    } if (EI_CLASSIFIER_RAW_SAMPLES_PER_FRAME != 3) {
ei_printf("ERR: EI_CLASSIFIER_RAW_SAMPLES_PER_FRAME should be equal to 3
(the 3 sensor axes)\n");
return;
    }
inference_thread.start(mbed::callback(&run_inference_background))
; }

/**
 * @brief Return the sign of the number

```

```

*
*   @param number
*   @return int 1 if positive (or 0) -1 if negative
*/ float ei_get_sign(float number)
{
    return (number >= 0.0) ? 1.0
: -1.0;
}

/**
*   @brief      Run inferencing in the background.
*/ void
run_inference_background()
{
    // wait until we have a full buffer    delay((EI_CLASSIFIER_INTERVAL_MS
* EI_CLASSIFIER_RAW_SAMPLE_COUNT) + 100);
    // This is a structure that smoothens the output result
    // With the default settings 70% of readings should be the same before
    classifying.
    ei_classifier_smooth_t smooth;
    ei_classifier_smooth_init(&smooth, 10 /* no. of readings */, 7 /*
min. readings the same */, 0.8 /* min. confidence */, 0.3 /* max
anomaly */);    while (1) {
        // copy the buffer                memcpy(inference_buffer, buffer,
EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE * sizeof(float));

        // Turn the raw buffer in a signal which we can the classify
        signal_t signal;
        int err = numpy::signal_from_buffer(inference_buffer,
EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE, &signal);        if (err != 0)
        {
            ei_printf("Failed to create signal from buffer
(%d)\n", err);                return;
        }

        // Run the classifier
        ei_impulse_result_t result = { 0 };
        err = run_classifier(&signal, &result, debug_nn);
        if (err != EI_IMPULSE_OK) {                ei_printf("ERR: Failed
to run classifier (%d)\n", err);                return;
        }

        // print the predictions                ei_printf("Predictions ");
        ei_printf("(DSP: %d ms., Classification: %d ms., Anomaly: %d ms.)",
            result.timing.dsp, result.timing.classification,
            result.timing.anomaly);                ei_printf(": ");

        // ei_classifier_smooth_update yields the predicted label
        const char *prediction = ei_classifier_smooth_update(&smooth, &result);
        ei_printf("%s ", prediction);                // print the cumulative results
        ei_printf(" [ ");

```



```

        for (size_t ix = 0; ix < smooth.count_size;
ix++) {
            ei_printf("%u", smooth.count[ix]);
if (ix != smooth.count_size + 1) {
ei_printf(", ");
}
else {
            ei_printf(" ");
}
ei_printf("]\n");

delay(run_inference_every_ms);
}
ei_classifier_smooth_free(&smooth
);
}

/**
 * @brief Get data and run inferencing
 *
 * @param[in] debug Get debug info if true
 */
void
loop()
{
    while (1)
    {
        // Determine the next tick (and then sleep later)
uint64_t next_tick = micros() + (EI_CLASSIFIER_INTERVAL_MS * 1000);
        // roll the buffer -3 points so we can overwrite the last one
numpy::roll(buffer, EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE, -3);

        // read to the end of the buffer
IMU.readAcceleration(
buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3],
buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 2],
buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 1]
);
        for (int i =
0; i < 3; i++) {
            if (fabs(buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3 +
i]) > MAX_ACCEPTED_RANGE) {
buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3 + i] =
ei_get_sign(buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3 + i]) *
MAX_ACCEPTED_RANGE;
}
}
}

```

```

        buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 3] *= CONVERT_G_TO_MS2;
        buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 2] *= CONVERT_G_TO_MS2;
        buffer[EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE - 1] *= CONVERT_G_TO_MS2;

        // and wait for next tick          uint64_t
time_to_wait = next_tick - micros();
delay((int)floor((float)time_to_wait / 1000.0f));
delayMicroseconds(time_to_wait % 1000);
    }
}

#ifdef EI_CLASSIFIER_SENSOR || EI_CLASSIFIER_SENSOR !=
EI_CLASSIFIER_SENSOR_ACCELEROMETER
#error "Invalid model for current sensor"
#e
nd
if

```

6) Screen shot of Arduino Terminal – Result

```

15:43:05.004 -> Predictions (DSP: 104 ms., Classification: 0 ms., Anomaly: 0 ms.): left-right [ 0, 0, 0, 10, 0, 0, ]
15:43:05.997 -> Predictions (DSP: 104 ms., Classification: 0 ms., Anomaly: 0 ms.): left-right [ 0, 0, 0, 10, 0, 0, ]
15:43:06.323 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): left-right [ 0, 0, 0, 10, 0, 0, ]
15:43:06.634 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): left-right [ 0, 0, 1, 9, 0, 0, ]
15:43:06.901 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): left-right [ 0, 0, 1, 8, 1, 0, ]
15:43:07.214 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): left-right [ 0, 0, 1, 7, 2, 0, ]
15:43:07.523 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): uncertain [ 0, 0, 1, 6, 3, 0, ]
15:43:07.848 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): uncertain [ 0, 0, 2, 5, 3, 0, ]
15:43:08.172 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): uncertain [ 0, 0, 2, 4, 4, 0, ]
15:43:08.448 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): uncertain [ 0, 0, 3, 3, 4, 0, ]

```

```

15:43:17.115 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 10, 0, 0, 0, ]
15:43:18.041 -> Predictions (DSP: 104 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 10, 0, 0, 0, ]
15:43:18.373 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 10, 0, 0, 0, ]
15:43:18.686 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 10, 0, 0, 0, ]
15:43:18.970 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 10, 0, 0, 0, ]
15:43:19.314 -> Predictions (DSP: 104 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 10, 0, 0, 0, ]
15:43:19.589 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 10, 0, 0, 0, ]
15:43:19.936 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 10, 0, 0, 0, ]
15:43:20.206 -> Predictions (DSP: 104 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 10, 0, 0, 0, ]
15:43:20.542 -> Predictions (DSP: 106 ms., Classification: 0 ms., Anomaly: 0 ms.): idle [ 0, 0, 9, 0, 1, 0, ]

```